Introduction: We are seeking a highly skilled Senior ML Engineer - Gen AI to join our innovative team. This role is pivotal in driving the development and deployment of cutting-edge generative AI solutions that transform business processes and customer experiences. The successful candidate will leverage advanced machine learning techniques and state-of-the-art tools to build scalable, intelligent systems that deliver significant impact across the organization.

Roles and Responsibilities:

Design, develop, and optimize machine learning models, focusing on generative AI and large language models (LLM) to solve complex business problems.
Implement and maintain robust MLOps pipelines for continuous integration, deployment, and monitoring of AI models in production environments.
Develop and integrate Agentic AI systems that enable autonomous decision-making and enhance AI-driven workflows.
Utilize vector search technologies to improve information retrieval and recommendation systems.
Build and manage applications using Lang Chain and Lang Graph frameworks to enhance natural language understanding and generation capabilities.
Collaborate with cross-functional teams including data scientists, software engineers, and product managers to translate business requirements into scalable AI solutions.
Ensure best practices in code quality, testing, and documentation are followed.
Optimize cloud infrastructure and AI workloads on AWS to ensure high availability, scalability, and cost-efficiency.
Stay abreast of emerging trends and advancements in generative AI and machine learning to continuously innovate and improve solutions.

Qualifications:

Bachelor’s or Master’s degree in Computer Science, Engineering, Data Science, or a related field. Advanced degrees preferred.
5-8 years of professional experience in machine learning engineering, with a strong focus on generative AI and large language models.
Proven expertise in Python programming and relevant ML libraries and frameworks.
Hands-on experience with MLOps tools and practices for model deployment and lifecycle management.
Deep understanding of LLM architectures and practical experience working with models like GPT, BERT, or similar.
Experience developing and deploying Agentic AI systems or autonomous AI agents.
Proficiency in implementing vector search algorithms and technologies such as FAISS or similar.
Practical knowledge of Lang Chain and Lang Graph for advanced language model applications.
Strong experience with cloud platforms, particularly AWS, including services related to AI/ML like SageMaker, Lambda, and EC2.
Excellent problem-solving skills, communication abilities, and teamwork orientation.

Tools and Technologies:

Machine Learning Frameworks: TensorFlow, PyTorch, Scikit-learn
Programming Languages: Python (primary)
MLOps: MLflow, Kubeflow, Airflow, Docker, Kubernetes
Large Language Models: GPT, BERT, OpenAI APIs
Agentic AI: Autonomous agent frameworks and custom implementations
Vector Search: FAISS, Annoy, Elasticsearch vector capabilities
Lang Chain & Lang Graph frameworks
Cloud Platforms: AWS (SageMaker, Lambda, EC2, S3)
Version control and collaboration: Git, GitHub/GitLab

Join us to be at the forefront of generative AI innovation, driving impactful solutions that redefine the future of intelligent systems.

skills:
Required Skills
Python
NumPy
Pandas
Scikit-learn
TensorFlow
PyTorch
Deep Learning
NLP
Data Preprocessing
Feature Engineering
Git
Preferred Skills
Computer Vision
Hugging Face
Docker
AWS
Responsibilities
Develop machine learning models.
Train and evaluate models.
Process datasets.
Deploy ML models.